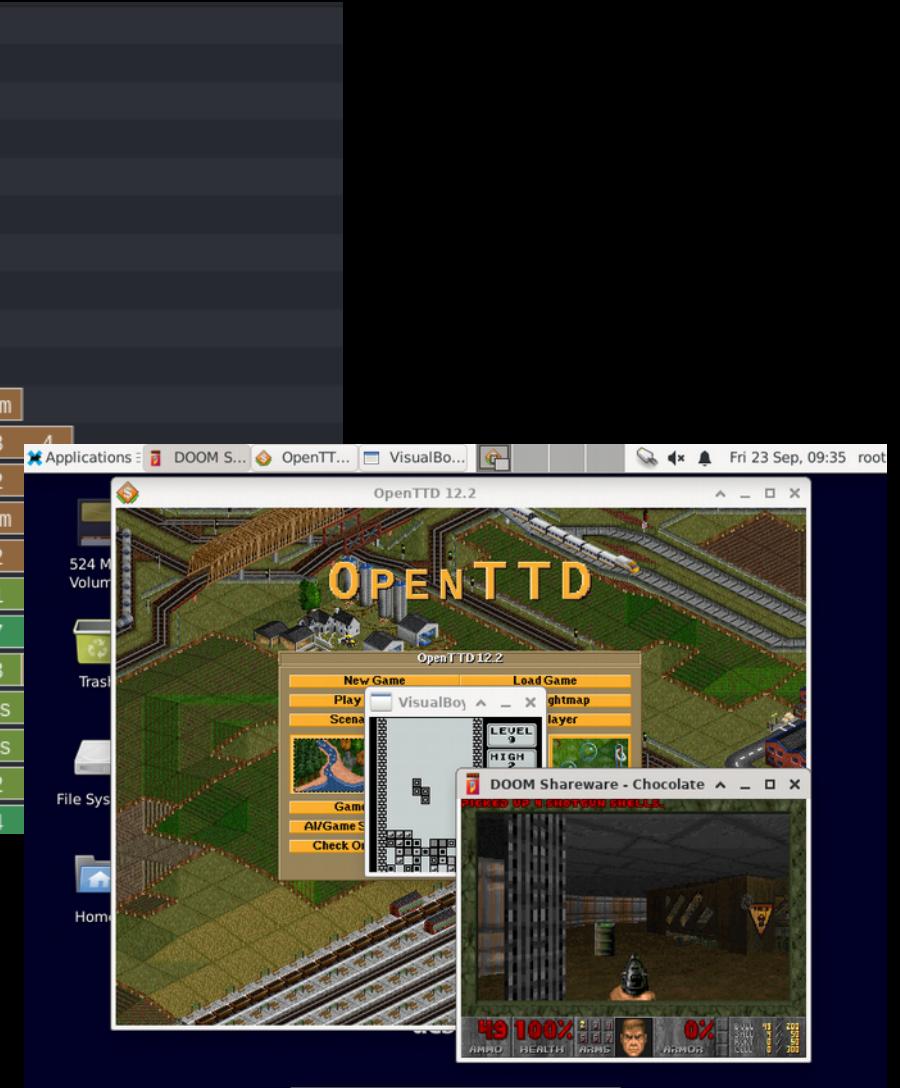


# Open CPU / SoC design, all the way up to Debian

lw	ra, 12(sp)	F 1 2   Rn 1   Ds 1   Is 1 2 3   Cm 1 2
andi	a5, a2, 255	F 1 2   Rn 1   Ds 1 2 3   Is Cm 1 2 3
sw	a5, 20(sp)	F 1 2   Rn 1   Ds 1   Is 1 2 3 4   Cm 1
bltz	ra, pc + 912	F 1 2   Rn 1   Ds 1 2 3 4 5   Is 1   Cm
beqz	s0, pc + 2064	F 1 2   Rn 1   Ds 1   Is 1   Cm 1 2 3 4
mv	t1, s0	F 1 2   Rn 1   Ds 1 2   Is   Cm 1 2 3 4
j	pc + 0xc	F 1 2   Rn 1   Ds 1   Is 1   Cm 1 2 3 4
lw	t2, 4(t1)	F 1 2   Rn 1   Ds 1 2   Is 1 2 3   Cm 1
lw	s3, 12(sp)	F 1 2   Rn 1   Ds 1   Is 1 2 3   Cm 1 2
lh	s2, 2(t2)	F 1 2   Rn 1   Ds 1 2 3 4   Is 1 2 3   Cm
bne	s2, s3, pc + 4076	F 1 2   Rn 1   Ds 1 2 3 4 5 6 7   Is 1   Cm
lw	t1, 0(t1)	F 1 2   Rn 1   Ds 1   Is 1 2 3   Cm 1 2 3
beqz	t1, pc + 20	F 1 2   Rn 1   Ds 1 2 3 4   Is 1   Cm 1 2
lw	t2, 4(t1)	F 1 2   Rn 1   Ds 1 2 3   Is 1 2 3   Cm
lw	s3, 12(sp)	F 1 2   Rn 1   Ds 1   Is 1 2 3   Cm 1 2
lh	s2, 2(t2)	F 1 2   Rn 1   Ds 1 2 3 4 5   Is 1
bne	s2, s3, pc + 4076	F 1 2   Rn 1   Ds 1 2 3 4 5 6 7
lw	t1, 0(t1)	F 1 2   Rn 1   Ds 1 2   Is 1 2 3
beqz	t1, pc + 20	F 1 2   Rn 1   Ds 1 2 3 4 5   Is
lw	t2, 4(t1)	F 1 2   Rn 1   Ds 1 2 3 4   Is
lw	s3, 12(sp)	F 1 2   Rn 1   Ds 1 2   Is 1 2
lh	s2, 2(t2)	F 1 2   Rn 1   Ds 1 2 3 4



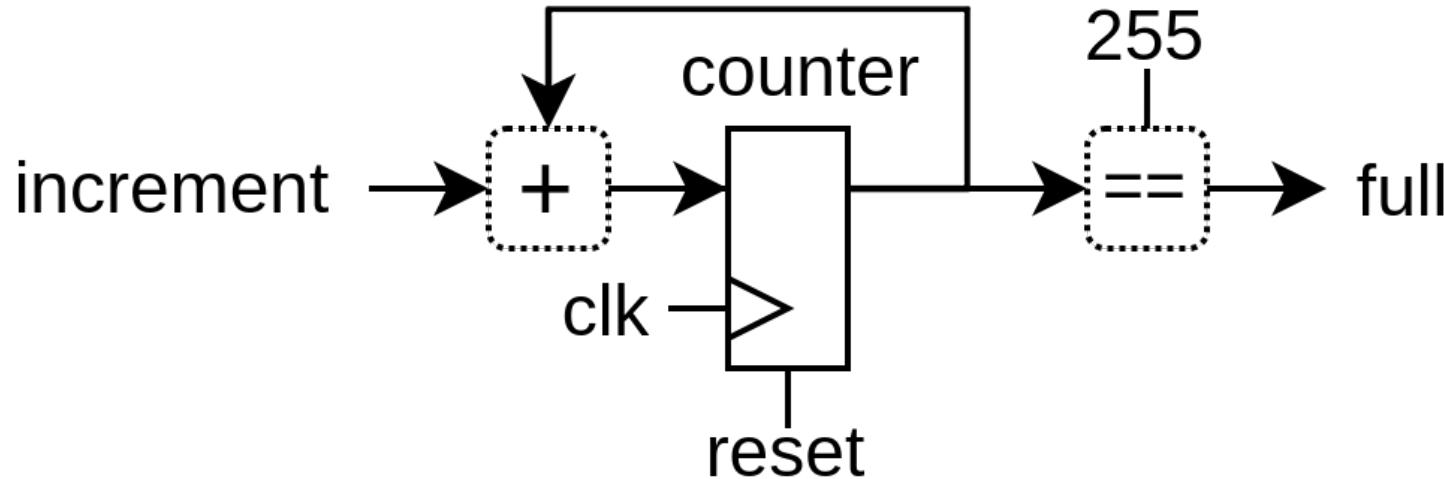
# Background / whoami

- Dolu1990 on github
- Active on open/free project
  - SpinalHDL / VexRiscv / NaxRiscv
- Software / Hardware background
  - Without computer science degree
  - With Industrial system / Electronic degree

# Introduction

- This talk will be based on NaxRiscv
  - <https://github.com/SpinalHDL/NaxRiscv>
- What the talk will do
  - Share experience / flow
  - Give tips / keywords

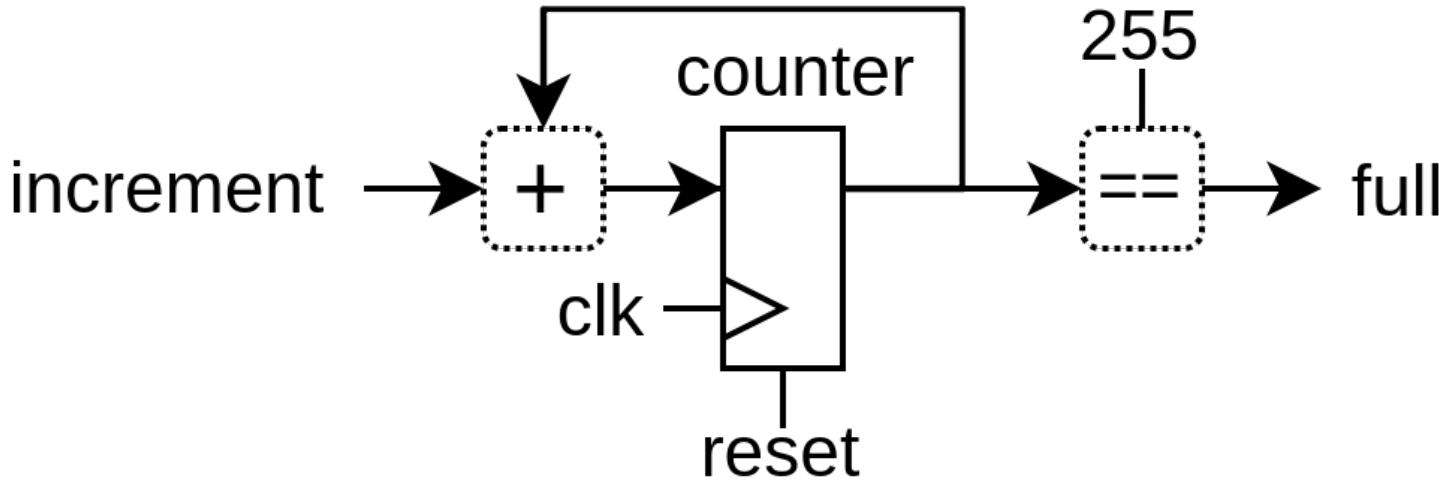
# Hardware description



- “There is no alternative” (VHDL / System-Verilog)
  - Scala based : SpinalHDL, Chisel
  - Python based : Migen, Amaranth
  - ...



# It can be pretty / explicit

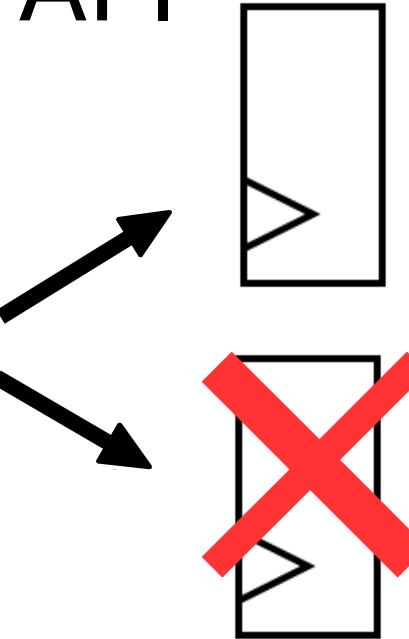


```
class Timer extends Component {
    val increment = in Bool()
    val counter = Reg(UInt(8 bits)) init(0)
    val full = counter === 255

    when(increment) {
        counter := counter + 1
    }
}
```

# Hardware description API

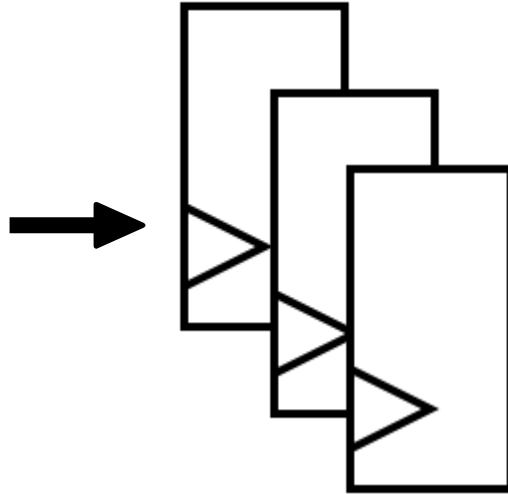
```
if(featureEnabled) {  
    Reg(UInt(8 bits))  
}
```



- Control flow : **if** / for
- Data structures : dynamic array / hash map / hash set
- Lambda function : reduce / fold / map / filter / ...
- OOP : class / software interface
- See <https://spinalhdl.github.io/NaxRiscv-Rtd/main/NaxRiscv/abstraction/index.html>

# Hardware description API

```
for(i <- 0 to 2){  
    Reg(UInt(8 bits))  
}
```



- Control flow : if / **for**
- Data structures : dynamic array / hash map / hash set
- Lambda function : reduce / fold / map / filter / ...
- OOP : class / software interface
- See <https://spinalhdl.github.io/NaxRiscv-Rtd/main/NaxRiscv/abstraction/index.html>

# Hardware description API

```
val flushes = ArrayBuffer[Flush] ()  
flushes += new Flush()  
flushes += new Flush()  
flushes += new Flush()
```



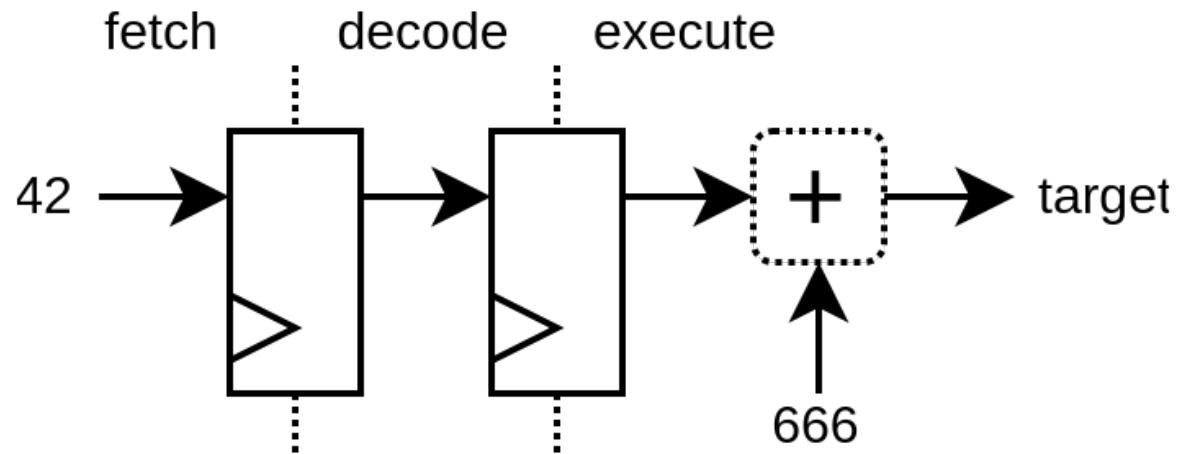
Index	Value
0	hardware
1	hardware
2	hardware

- Data structures : **dynamic array** / hash map / hash set
- Lambda function : reduce / fold / map / filter / ...
- OOP : class / software interface
- See <https://spinalhdl.github.io/NaxRiscv-Rtd/main/NaxRiscv/abstraction/index.html>

```

val PC = Hardtype(UInt(32 bits))
val instruction = Hardtype(UInt(32 bits))
val fetch, decode, execute = HashMap[Hardtype, Hardware]()
...
fetch(PC) = U(42)
...
val target = execute(PC) + U(666)
...
autoPipeline(fetch, decode, execute)

```



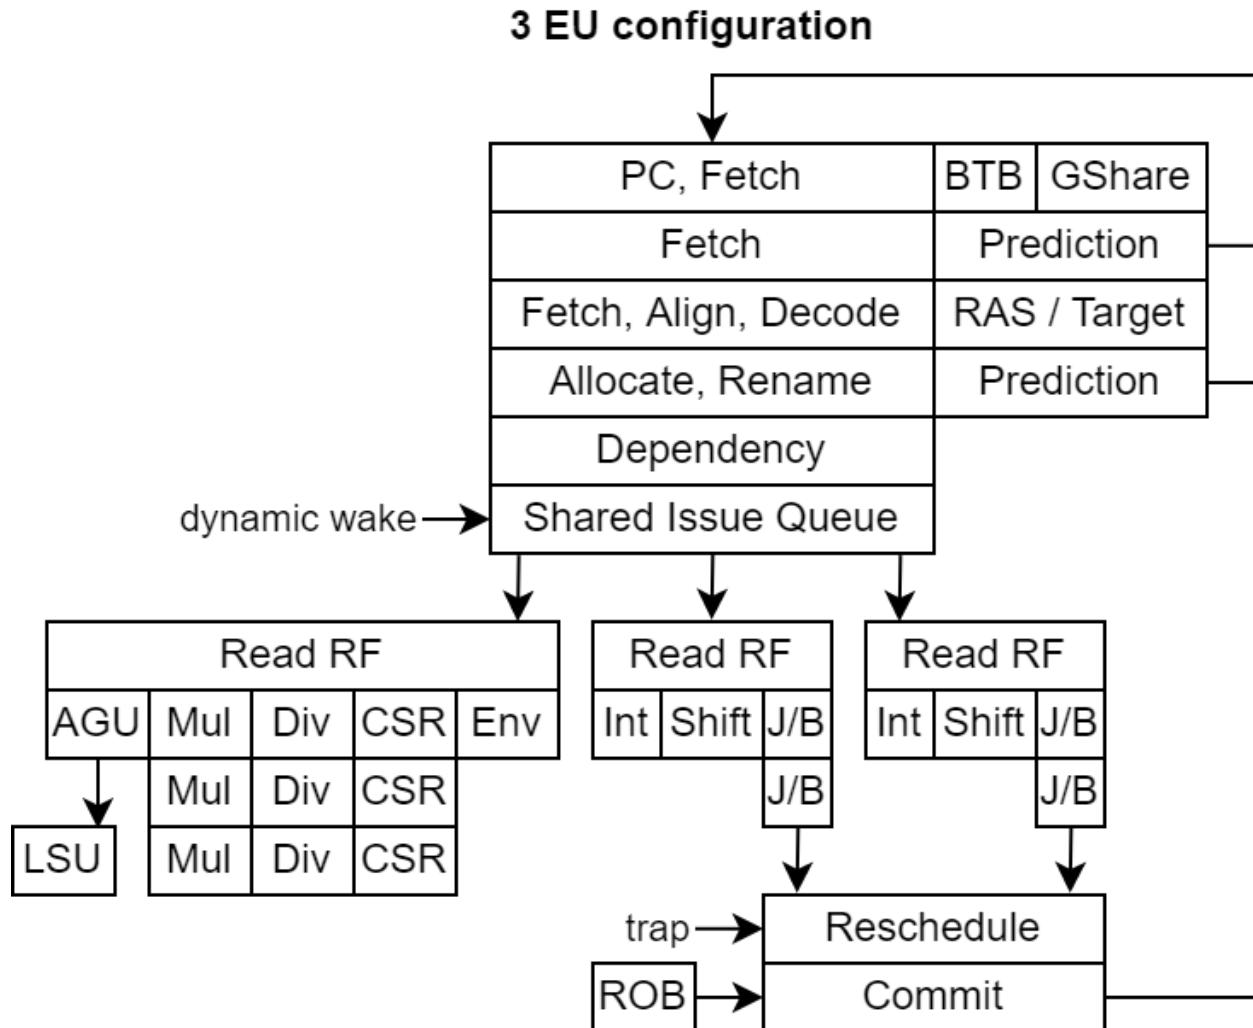
- Data structures : dynamic array / **hash map** / hash set
- Lambda function : reduce / fold / map / filter / ...
- OOP : class / software interface
- See <https://spinalhdl.github.io/NaxRiscv-Rtd/main/NaxRiscv/abstraction/index.html>

# ISA (instruction set architecture)

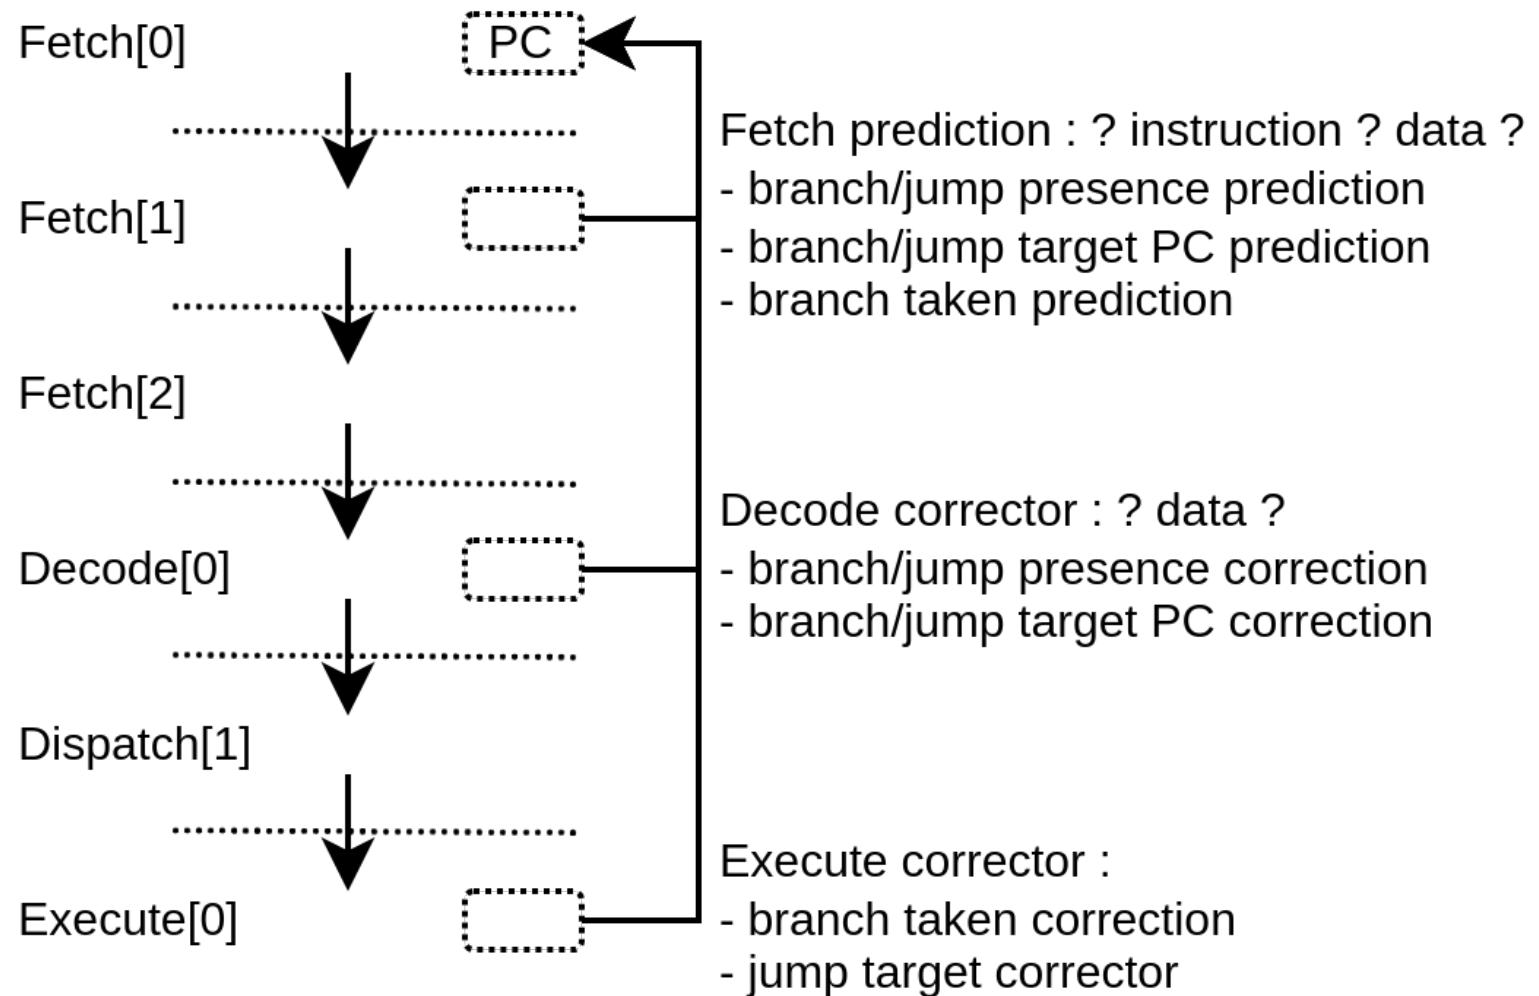


- A few open / free ISA (OpenRISC, **RISC-V**, ...)
- RISC-V
  - <https://riscv.org/technical/specifications/>
  - Secretly being embedded in ASICs
  - Mostly bloat free, from a FPGA perspective)
  - GCC / LLVM / Qemu
  - Linux / Debian port
  - ...

# CPU design (NaxRiscv)



# “Branch prediction”





# Secret leak (Konata trace)

jal	pc + 0x230	1	F 1 2   Rn 1   Ds 1   Is 1   Cm 1 2 3 4 5 6 7 8 9 10 11 12 13
addi	sp, sp, 4088		F 1 2   Rn 1   Ds 1   Is   Cm 1 2 3 4 5 6 7 8 9 10 11 12 13 14
sd	a3, 0(sp)		F 1 2   Rn 1   Ds 1   Is 1 2 3 4   Cm 1 2 3 4 5 6 7 8 9 10
ld	a3, 0(sp)	2	F 1 2   Rn 1   Ds 1 2   Is 1 2 3 4 5 6 7 8 9   Cm 1 2 3 4
sd	a3, 0(sp)		F 1 2   Rn 1   Ds 1 2   Is 1 2 3 4 5 6 7 8 9 10 11   Cm 1 2
ld	a3, 0(sp)		F 1 2   Rn 1   Ds 1 2 3   Is 1 2 3 4 5 6 7 8 9 10 11 12 13 14   Cm
beqz	a3, pc + 28	3	F 1 2   Rn 1   Ds 1 2 3 4 5 6 7 8 9 10 11 12 13 14 15 16 17 18 19   Is 1   Cm
lb	t2, 0(a0)	4	F 1 2   Rn 1   Ds 1 2 3
srl	t2, t2, a1		F 1 2   Rn 1   Ds 1 2 3 4 5   Is
andi	t2, t2, 1		F 1 2   Rn 1   Ds 1 2 3 4 5 6   Is
slli	t2, t2, 6		F 1 2   Rn 1   Ds 1 2 3 4 5 6   Is
or	t2, t2, a2		F 1 2   Rn 1   Ds 1 2 3 4 5 6 7   Is
ld	t2, 0(t2)	5	F 1 2   Rn 1   Ds 1 2 3 4 5 6 7   Is 1 2 3
addi	sp, sp, 8		F 1 2   Rn 1   Ds 1   Is
ret			F 1 2   Rn 1   Ds 1   Is 1
mv	a2, s1		F 1 2
nop			F 1 2   Rn 1   Ds 1 2   Is
nop			F 1 2   Rn 1   Ds 1   Is
nop			F 1 2   Rn 1   Ds 1 2   Is
nop			F 1 2   Rn 1   Ds 1 2   Is
nop			F 1 2   Rn 1   Ds 1 2   Is
nop			F 1 2   Rn 1   Ds 1 2   Is
nop			F 1 2
nop			F 1 2
addi	sp, sp, 8	6	F 1 2
ret			F 1 2

# Secret leak

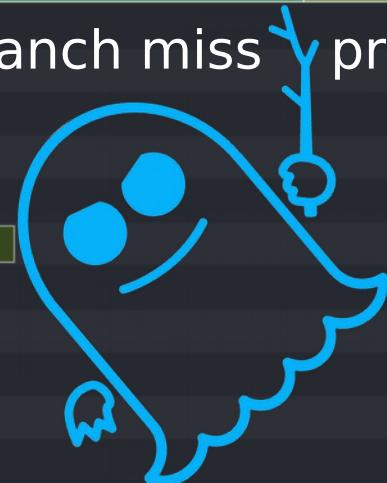
jal	pc + 0x230	1	F 1 2   Rn 1   Ds 1   Is 1   Cm 1 2 3 4 5 6 7 8 9 10 11 12 13
addi	sp, sp, 4088		F 1 2   Rn 1   Ds 1   Is   Cm 1 2 3 4 5 6 7 8 9 10 11 12 13 14
sd	a3, 0(sp)		F 1 2   Rn 1   Ds 1   Is 1 2 3 4   Cm 1 2 3 4 5 6 7 8 9 10
ld	a3, 0(sp)	2	F 1 2   Rn 1   Ds 1 2   Is 1 2 3 4 5 6 7 8 9   Cm 1 2 3 4
sd	a3, 0(sp)		F 1 2   Rn 1   Ds 1 2   Is 1 2 3 4 5 6 7 8 9 10 11   Cm 1 2
ld	a3, 0(sp)		F 1 2   Rn 1   Ds 1 2 3   Is 1 2 3 4 5 6 7 8 9 10 11 12 13 14   Cm
beqz	a3, pc + 28	3	F 1 2   Rn 1   Ds 1 2 3 4 5 6 7 8 9 10 11 12 13 14 15 16 17 18 19   Is 1   Cm
lb	t2, 0(a0)	4	F 1 2   Rn 1   Ds 1 2 3
srl	t2, t2, a1		F 1 2   Rn 1   Ds 1 2 3 4   Is
andi	t2, t2, 1		F 1 2   Rn 1   Ds 1 2 3 4 6   Is
slli	t2, t2, 6		F 1 2   Rn 1   Ds 1 2 3 4 5 6   Is
or	t2, t2, a2		F 1 2   Rn 1   Ds 1 2 3 4 5 6 7   Is
ld	t2, 0(t2)	5	F 1 2   Rn 1   Ds 1 2 3 4 5 6 7   Is 1 2 3
addi	sp, sp, 8		F 1 2   Rn 1   Ds 1   Is
ret			F 1 2   Rn 1   Ds 1 2   Is
mv	a2, s1		F 1 2   Rn 1   Ds 1 2   Is
nop			F 1 2   Rn 1   Ds 1 2   Is
nop			F 1 2   Rn 1   Ds 1   Is
nop			F 1 2   Rn 1   Ds 1 2   Is
nop			F 1 2   Rn 1   Ds 1 2   Is
nop			F 1 2   Rn 1   Ds 1 2   Is
nop			F 1 2   Rn 1   Ds 1 2   Is
addi	sp, sp, 8	6	F 1 2   Rn 1   Ds 1 2   Is
ret			F 1 2   Rn 1   Ds 1 2   Is

Generate long dependencies

# Secret leak

jal	pc + 0x230	1	F 1 2   Rn 1   Ds 1   Is 1   Cm 1 2 3 4 5 6 7 8 9 10 11 12 13
addi	sp, sp, 4088		F 1 2   Rn 1   Ds 1   Is   Cm 1 2 3 4 5 6 7 8 9 10 11 12 13 14
sd	a3, 0(sp)		F 1 2   Rn 1   Ds 1   Is 1 2 3 4   Cm 1 2 3 4 5 6 7 8 9 10
ld	a3, 0(sp)	2	F 1 2   Rn 1   Ds 1 2   Is 1 2 3 4 5 6 7 8 9   Cm 1 2 3 4
sd	a3, 0(sp)		F 1 2   Rn 1   Ds 1 2   Is 1 2 3 4 5 6 7 8 9 10 11   Cm 1 2
ld	a3, 0(sp)		F 1 2   Rn 1   Ds 1 2 3   Is 1 2 3 4 5 6 7 8 9 10 11 12 13 14   Cm
beqz	a3, pc + 28	3	F 1 2   Rn 1   Ds 1 2 3 4 5 6 7 8 9 10 11 12 13 14 15 16 17 18 19   Is 1   Cm
lb	t2, 0(a0)	4	F 1 2   Rn 1   Ds 1 2 3
srl	t2, t2, a1		F 1 2   Rn 1   Ds 1 2 3 4 5   Is
andi	t2, t2, 1		F 1 2   Rn 1   Ds 1 2 3 4 5 6   Is
slli	t2, t2, 6		F 1 2   Rn 1   Ds 1 2 3 4 5 6   Is
or	t2, t2, a2		F 1 2   Rn 1   Ds 1 2 3 4 5 6 7   Is
ld	t2, 0(t2)	5	F 1 2   Rn 1   Ds 1 2 3 4 5 6 7   Is 1 2 3
addi	sp, sp, 8		F 1 2   Rn 1   Ds 1   Is
ret			F 1 2   Rn 1   Ds 1   Is 1
mv	a2, s1		F 1 2
nop			F 1 2   Rn 1   Ds 1 2   Is
nop			F 1 2   Rn 1   Ds 1   Is
nop			F 1 2   Rn 1   Ds 1 2   Is
nop			F 1 2   Rn 1   Ds 1 2   Is
nop			F 1 2   Rn 1   Ds 1 2   Is
nop			F 1 2   Rn 1   Ds 1 2   Is
addi	sp, sp, 8	6	F 1 2
ret			F 1 2

Laaaaaaaate branch miss prediction



# Secret leak

jal	pc + 0x230	1	F 1 2   Rn 1   Ds 1   Is 1   Cm 1 2 3 4 5 6 7 8 9 10 11 12 13
addi	sp, sp, 4088		F 1 2   Rn 1   Ds 1   Is   Cm 1 2 3 4 5 6 7 8 9 10 11 12 13 14
sd	a3, 0(sp)		F 1 2   Rn 1   Ds 1   Is 1 2 3 4   Cm 1 2 3 4 5 6 7 8 9 10
ld	a3, 0(sp)	2	F 1 2   Rn 1   Ds 1 2   Is 1 2 3 4 5 6 7 8 9   Cm 1 2 3 4
sd	a3, 0(sp)		F 1 2   Rn 1   Ds 1 2   Is 1 2 3 4 5 6 7 8 9 10 11   Cm 1 2
ld	a3, 0(sp)		F 1 2   Rn 1   Ds 1 2 3   Is 1 2 3 4 5 6 7 8 9 10 11 12 13 14   Cm
beqz	a3, pc + 28	3	F 1 2   Rn 1   Ds 1 2 3 4 5 6 7 8 9 10 11 12 13 14 15 16 17 18 19   Is 1   Cm
lb	t2, 0(a0)	4	F 1 2   Rn 1   Ds 1 2 3
srl	t2, t2, a1		F 1 2   Rn 1   Ds 1 2 3 4 5   Is
andi	t2, t2, 1		F 1 2   Rn 1   Ds 1 2 3 4 5 6   Is
slli	t2, t2, 6		F 1 2   Rn 1   Ds 1 2 3 4 5 6   Is
or	t2, t2, a2		F 1 2   Rn 1   Ds 1 2 3 4 5 6 7   Is
ld	t2, 0(t2)	5	F 1 2   Rn 1   Ds 1 2 3 4 5 6 7   Is 1 2 3
addi	sp, sp, 8		F 1 2   Rn 1   Ds 1   Is 1
ret			F 1 2
mv	a2, s1		F 1 2   Rn 1   Ds 1 2   Is
nop			F 1 2   Rn 1   Ds 1 2   Is
nop			F 1 2   Rn 1   Ds 1 2   Is
nop			F 1 2   Rn 1   Ds 1 2   Is
nop			F 1 2   Rn 1   Ds 1 2   Is
nop			F 1 2   Rn 1   Ds 1 2   Is
nop			F 1 2   Rn 1   Ds 1 2   Is
addi	sp, sp, 8	6	F 1 2   Rn 1   Ds 1 2   Is
ret			F 1 2

RAM[secret] -> (page fault)



# Secret leak

jal	pc + 0x230	1	F 1 2   Rn 1   Ds 1   Is 1   Cm 1 2 3 4 5 6 7 8 9 10 11 12 13
addi	sp, sp, 4088		F 1 2   Rn 1   Ds 1   Is   Cm 1 2 3 4 5 6 7 8 9 10 11 12 13 14
sd	a3, 0(sp)		F 1 2   Rn 1   Ds 1   Is 1 2 3 4   Cm 1 2 3 4 5 6 7 8 9 10
ld	a3, 0(sp)	2	F 1 2   Rn 1   Ds 1 2   Is 1 2 3 4 5 6 7 8 9   Cm 1 2 3 4
sd	a3, 0(sp)		F 1 2   Rn 1   Ds 1 2   Is 1 2 3 4 5 6 7 8 9 10 11   Cm 1 2
ld	a3, 0(sp)		F 1 2   Rn 1   Ds 1 2 3   Is 1 2 3 4 5 6 7 8 9 10 11 12 13 14   Cm
beqz	a3, pc + 28	3	F 1 2   Rn 1   Ds 1 2 3 4 5 6 7 8 9 10 11 12 13 14 15 16 17 18 19   Is 1   Cm
lb	t2, 0(a0)	4	F 1 2   Rn 1   Ds 1 2 3
srl	t2, t2, a1		F 1 2   Rn 1   Ds 1 2 3 4 5   Is
andi	t2, t2, 1		F 1 2   Rn 1   Ds 1 2 3 4 5 6   Is
slli	t2, t2, 6		F 1 2   Rn 1   Ds 1 2 3 4 5 6   Is
or	t2, t2, a2		F 1 2   Rn 1   Ds 1 2 3 4 5 6   Is
ld	t2, 0(t2)	5	F 1 2   Rn 1   Ds 1 2 3 4 5 6 7   Is 1 2 3
addi	sp, sp, 8		F 1 2   Rn 1   Ds 1   Is
ret			F 1 2   Rn 1   Ds 1   Is 1
mv	a2, s1		F 1 2
nop			F 1 2   Rn 1   Ds 1 2   Is
nop			F 1 2   Rn 1   Ds 1   Is
nop			F 1 2   Rn 1   Ds 1 2   Is
nop			F 1 2   Rn 1   Ds 2   Is
nop			F 1 2   Rn 1   Ds 1 2   Is
nop			F 1 2   Rn 1   Ds 1 2   Is
addi	sp, sp, 8	6	F 1 2
ret			F 1 2

probe + RAM[secret]



# Secret leak

jal	pc + 0x230	1	F 1 2   Rn 1   Ds 1   Is 1   Cm 1 2 3 4 5 6 7 8 9 10 11 12 13
addi	sp, sp, 4088		F 1 2   Rn 1   Ds 1   Is   Cm 1 2 3 4 5 6 7 8 9 10 11 12 13 14
sd	a3, 0(sp)		F 1 2   Rn 1   Ds 1   Is 1 2 3 4   Cm 1 2 3 4 5 6 7 8 9 10
ld	a3, 0(sp)	2	F 1 2   Rn 1   Ds 1 2   Is 1 2 3 4 5 6 7 8 9   Cm 1 2 3 4
sd	a3, 0(sp)		F 1 2   Rn 1   Ds 1 2   Is 1 2 3 4 5 6 7 8 9 10 11   Cm 1 2
ld	a3, 0(sp)		F 1 2   Rn 1   Ds 1 2 3   Is 1 2 3 4 5 6 7 8 9 10 11 12 13 14   Cm
beqz	a3, pc + 28	3	F 1 2   Rn 1   Ds 1 2 3 4 5 6 7 8 9 10 11 12 13 14 15 16 17 18 19   Is 1   Cm
lb	t2, 0(a0)	4	F 1 2   Rn 1   Ds 1 2 3
srl	t2, t2, a1		F 1 2   Rn 1   Ds 1 2 3 4 5   Is
andi	t2, t2, 1		F 1 2   Rn 1   Ds 1 2 3 4 5 6   Is
slli	t2, t2, 6		F 1 2   Rn 1   Ds 1 2 3 4 5 6   Is
or	t2, t2, a2		F 1 2   Rn 1   Ds 1 2 3 4 5 6 7   Is
ld	t2, 0(t2)	5	F 1 2   Rn 1   Ds 1 2 3 4 5 6 7   Is 1 2 3
addi	sp, sp, 8		F 1 2   Rn 1   Ds 1   Is
ret			F 1 2   Rn 1   Ds 1   Is 1
mv	a2, s1		F 1 2
nop			F 1 2   Rn 1   Ds 1 2   Is
nop			F 1 2   Rn 1   Ds 1   Is
nop			F 1 2   Rn 1   Ds 1 2   Is
nop			F 1 2   Rn 1   Ds 1 2   Is
nop			F 1 2   Rn 1   Ds 1 2   Is
nop			F 1 2   Rn 1   Ds 1 2   Is
nop			F 1 2   Rn 1   Ds 1 2   Is
addi	sp, sp, 8	6	F 1 2   Rn 1   Ds 1 2   Is
ret			F 1 2

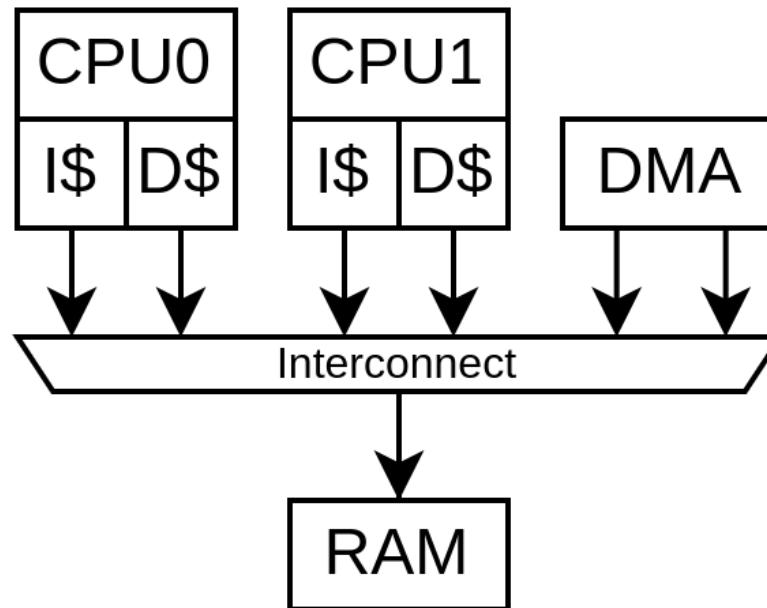
RAM[probe + RAM[secret]]



F 1 2 | Rn 1 | Ds 1 2 | Is

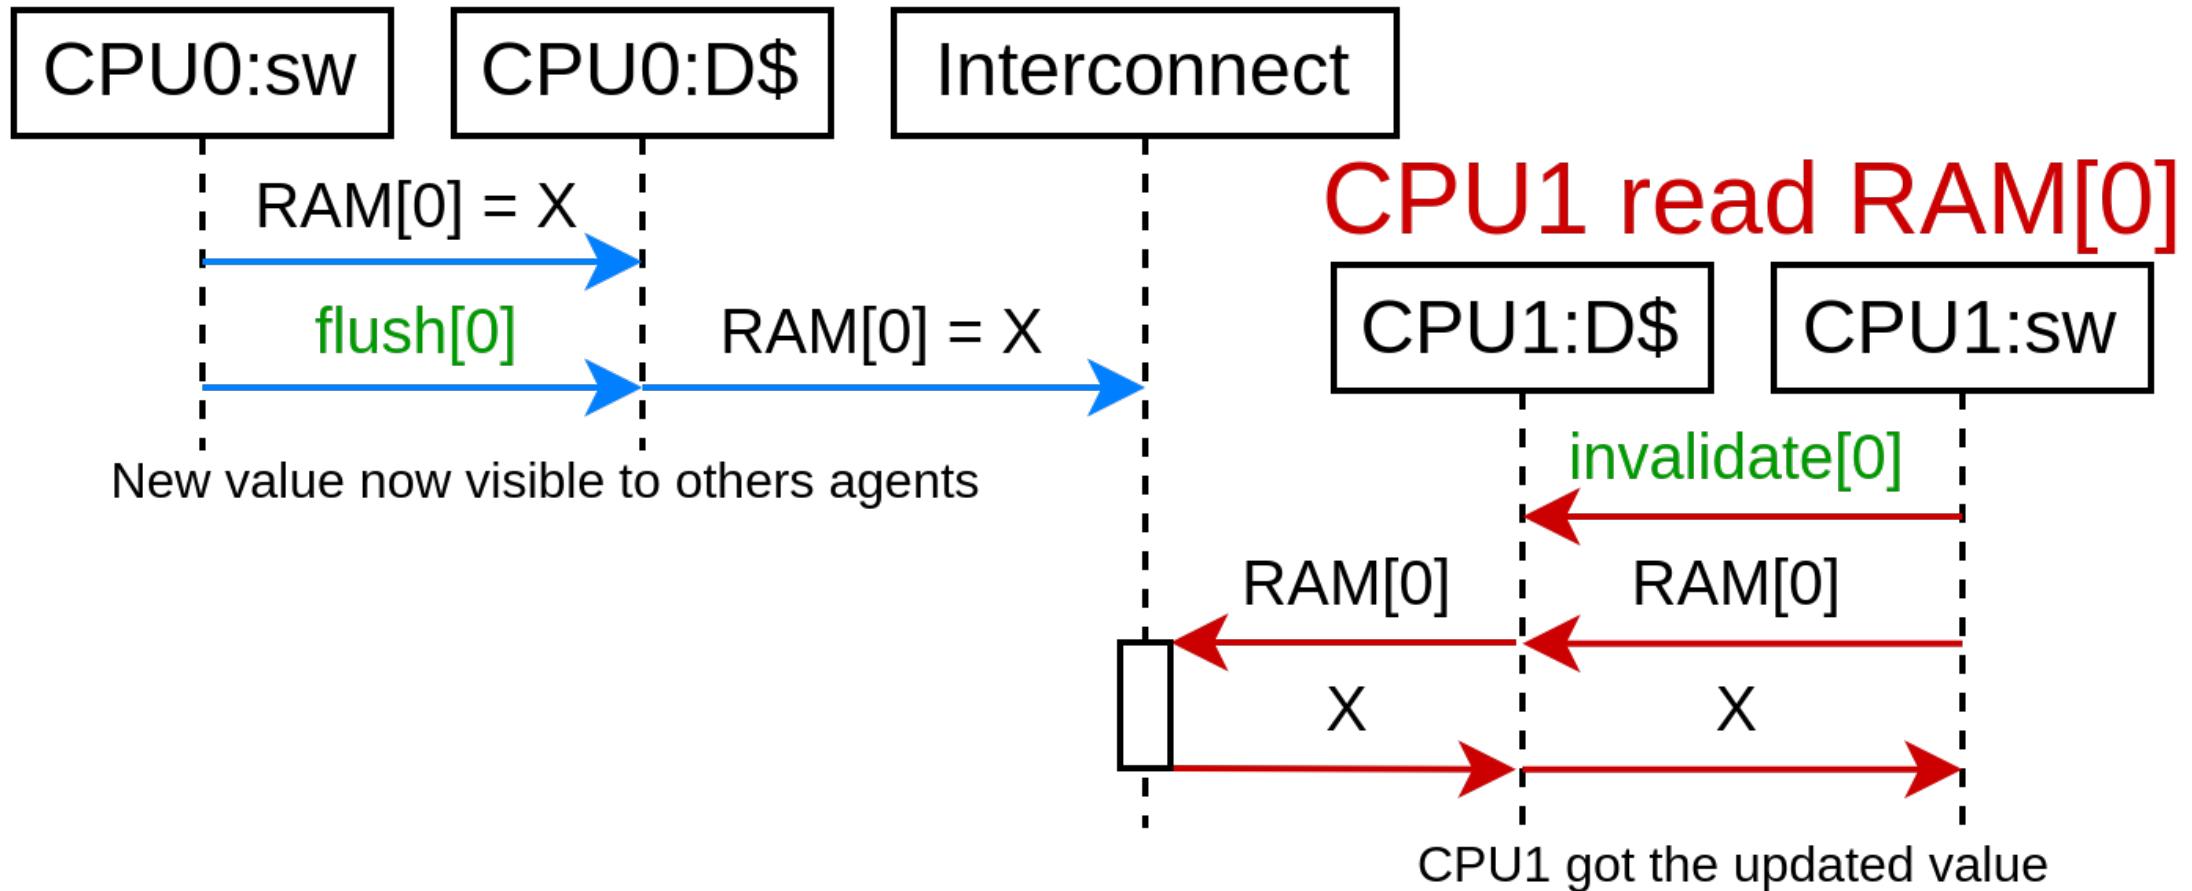
# Multi core requirements (RISC-V linux)

- Hardware cache coherency
- Inter-CPU software interrupts



# Software cache coherency

CPU0 write RAM[0]

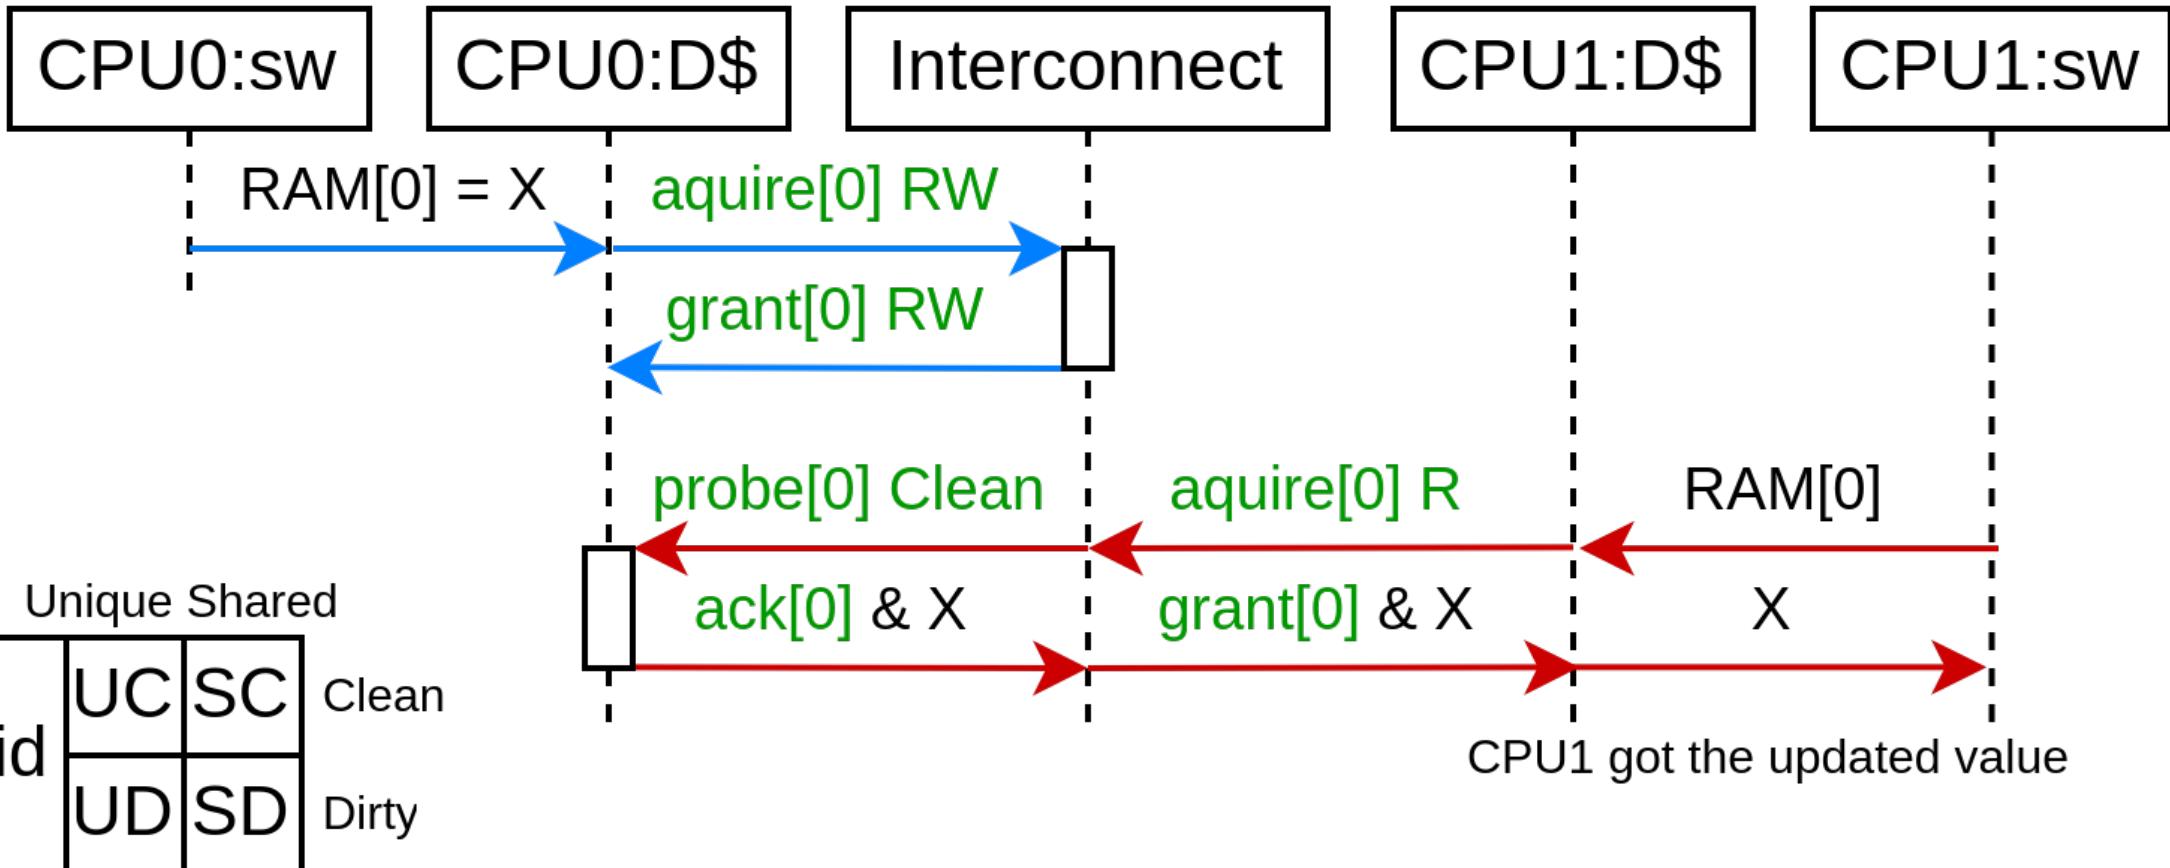


# Hardware cache coherency

with memory block copy and permissions

CPU0 write RAM[0]

CPU1 read RAM[0]



# Hardware cache coherency

And its side effects

```
volatile int array[2];

void thread_0() {
    while(1) {
        array[0] += 1;
    }
}

void thread_1() {
    while(1) {
        array[1] += 1;
    }
}
```

Same memory block (64 bytes): 18486 ps/iteration  
Different memory blocks : 902 ps/iteration

(on a 2 years old desktop CPU)

# Cache coherency

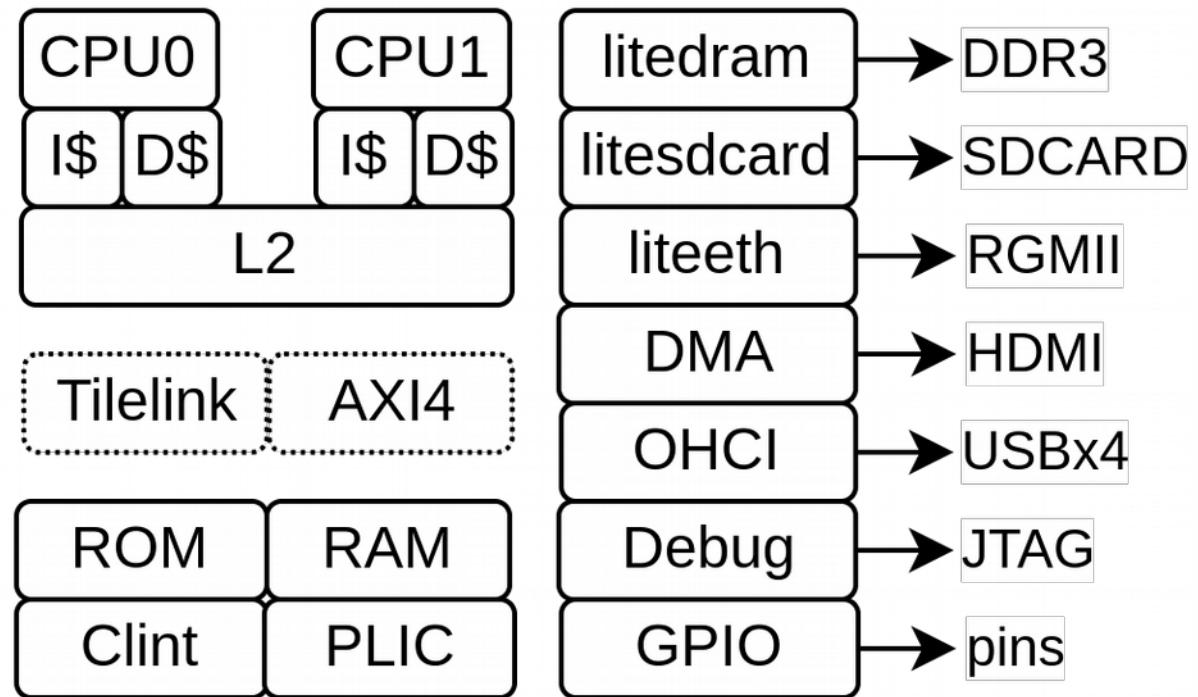
- Open specification : Tilelink
- Multiple open source implementations, ex :
  - <https://github.com/SpinalHDL/SpinalHDL/tree/dev/lib/src/main/scala/spinal/lib/bus/tilelink>
- Coherent L2 design
  - [https://spinalhdl.github.io/SpinalDoc-RTD/master/SpinalHDL/Libraries/Bus/tilelink\\_tilelink\\_fabric.html](https://spinalhdl.github.io/SpinalDoc-RTD/master/SpinalHDL/Libraries/Bus/tilelink_tilelink_fabric.html)
  - Exposes a lot of performance tricks
    - 1 pending transaction per 64 bytes block
    - 64 bytes per burst max
    - Keep bursts aligned
    - Alignment Aliasing

# Memory / Peripherals / Deployment

- LiteX
  - Open-source SoC framework (in Python)
  - Integrate many peripherals
    - DDRx controller
    - SDCARD
    - Ethernet, USB host
    - SATA, PCIe, ...

# Memory / Peripherals / Deployment

- `python3 -m litex_boards.targets.digilent_nexys_video --cpu-type=naxriscv --bus-standard axi-lite --with-video-framebuffer --with-coherent-dma --with-sdcard --with-ethernet --xlen=64 --scala-args='rvc=true,rvf=true,rvd=true,alu-count=2,decode-count=2' --with-jtag-tap --sys-clk-freq 100000000 --cpu-count 2 --build -load`
- + stuff for USB / JTAG



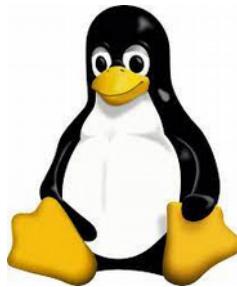
# USB host support

- USB host support =>
  - Keyboard / mouse / audio / flash drive / sdcard / uart / ethernet / bluetooth
- Open specification :
  - Open Host Controller Interface Specification for USB (OHCI)
  - USB 12 Mb/s, up to 15 ports, 2 FPGA GPIO per port
- Open implementation :
  - [https://spinalhdl.github.io/SpinalDoc-RTD/master/SpinalHDL/Libraries/Com/usb\\_ohci.html](https://spinalhdl.github.io/SpinalDoc-RTD/master/SpinalHDL/Libraries/Com/usb_ohci.html)
  - <https://github.com/SpinalHDL/SpinalHDL/blob/dev/lib/src/main/scala/spinal/lib/com/usb/ohci/UsbOhci.scala>
  - [https://github.com/litex-hub/pythondata-misc-usb\\_ohci/tree/master/pythondata\\_misc\\_usb\\_ohci/verilog](https://github.com/litex-hub/pythondata-misc-usb_ohci/tree/master/pythondata_misc_usb_ohci/verilog)

# Synthesis / Place / Route

- Some open source tools for
  - Xilinx 7-Series
  - Lattice ice40 / ecp5
  - ...
- Synthesis : Yosys
- Place and route : Nextpnr
- ...

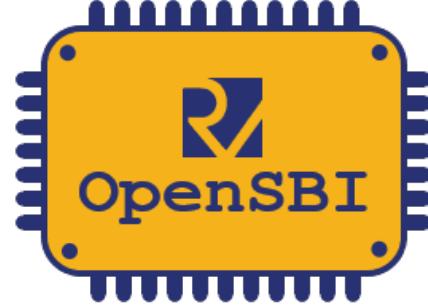
# Running Linux



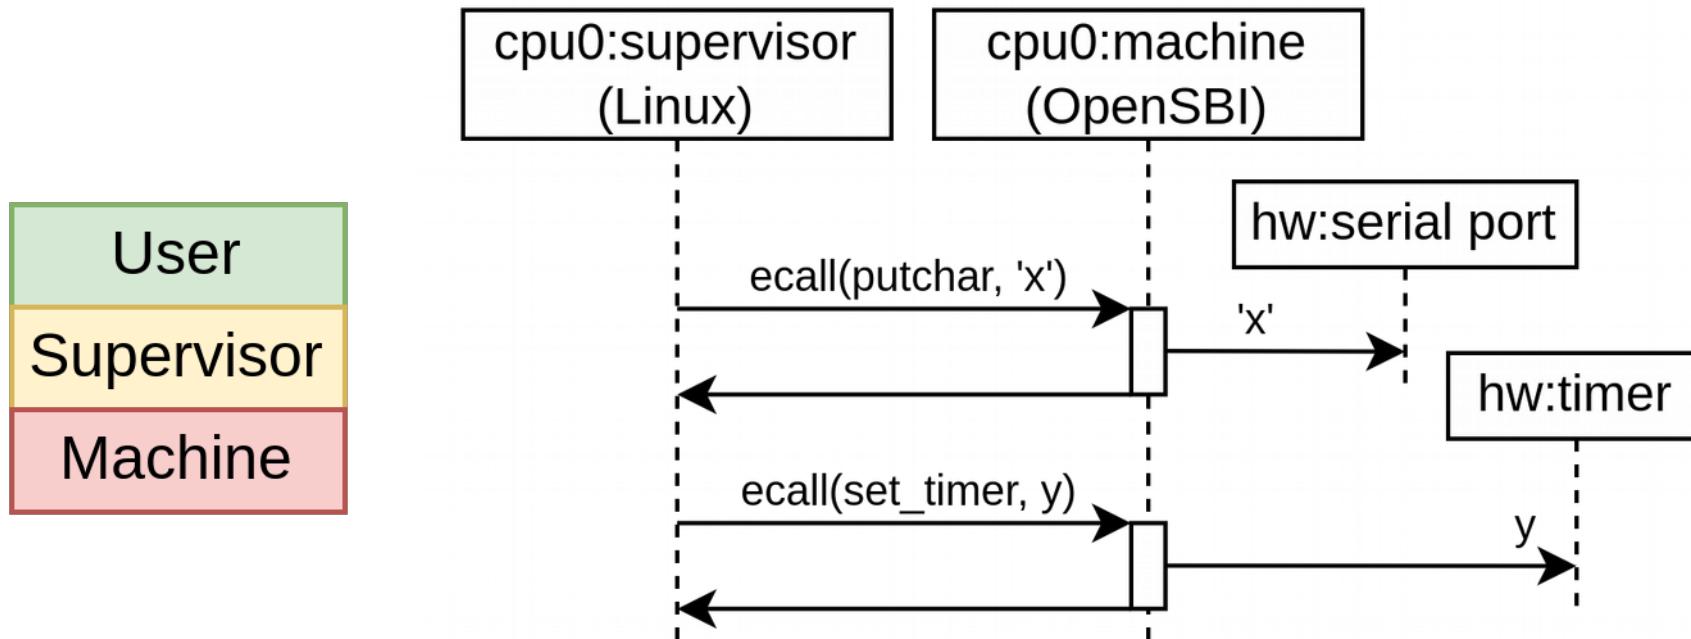
- Requirements on RISC-V
  - Hardware cache coherency (between CPUs + DMA)
  - Machine/Supervisor/User modes (riscv-privileged spec)
  - MMU / Atomic
  - SBI (Supervisor Binary Interface, kinda like BIOS)



# SBI (Supervisor Binary Interface)



- <https://github.com/riscv-non-isa/riscv-sbi-doc>
- <https://github.com/riscv-software-src/opensbi>



# Running Debian

- Requirements on RISC-V
  - RV64IMAFDC
  - Linux kernel
  - Some storage (ex: SDCARD)
- Doc : [https://github.com/SpinalHDL/NaxSoftware/tree/main/debian\\_litex](https://github.com/SpinalHDL/NaxSoftware/tree/main/debian_litex)
- Video : <https://x.com/dolu1990/status/1712848731349889108>

```
0[*]                                     0.6%] Tasks: 28, 18 thr, 61 kthr; 1 running
1[#####]                                13.8%] Load average: 3.48 1.37 0.50
Mem[|||||##*@@@@@@]                      46.2M/472M] Uptime: 00:01:43
Swp[                                         0K/4.00G]

[Main] [I/O]
 PID USER      PRI  NI   VIRT   RES   SHR   S CPU%vMEM%  TIME+  Command
 512 root      20   0  3784  2776  2220  R 11.6   0.6  0:01.42 htop
 405 root      20   0 25000  7160  5568  S  0.6   1.5  0:00.89 /usr/sbin/cup
    1 root      20   0 16768 10876  7524  S  0.0   2.2  0:28.04 /sbin/init
```



# Verification / Debugging

- May create rituals
- May exhaust your sanity
- May never end
- May crush your soul

# Verification / Debugging

- May create rituals
- May exhaust your sanity
- May never end
- May crush your soul

```
[1377614.579833] rcu: INFO: rcu_sched self-detected stall on CPU
[1377614.579845] rcu: 18-....: (2099 ticks this GP) idle=54e/1/0x4000000000000002
[1377614.579895] (t=2100 jiffies g=155867385 q=20879)
[1377614.579898] Task dump for CPU 18:
[1377614.579899] CPU 1/KVM R running task 0 1030947 256019 0x06000004
[1377614.579902] Call Trace:
[1377614.579912] ([<0000001f1f4b4f52>] show_stack+0x7a/0xc0)
[1377614.579918] [<0000001f1ec8e96c>] sched_show_task.part.0+0xdc/0x100
[1377614.579919] [<0000001f1f4b7248>] rcu_dump_cpu_stacks+0xc0/0x100
[1377614.579924] [<0000001f1ecdd10c>] rcu_sched_clock_irq+0x75c/0x980
[1377614.579926] [<0000001f1eceb26c>] update_process_times+0x3c/0x80
[1377614.579931] [<0000001f1ecfcfea>] tick_sched_handle.isra.0+0x4a/0x70
[1377614.579932] [<0000001f1ecfd28e>] tick_sched_timer+0x5e/0xc0
[1377614.579933] [<0000001f1ecec294>] __hrtimer_run_queues+0x114/0x2f0
[1377614.579935] [<0000001f1ecefcdc>] hrtimer_interrupt+0x12c/0x2a0
[1377614.579938] [<0000001f1ebecb6a>] do_IRQ+0xaa/0xb0
[1377614.579942] [<0000001f1f4c6d08>] ext_int_handler+0x130/0x134
[1377614.579945] [<0000001f1ec0af10>] ptep_zap_key+0x40/0x60
```

Hello darkness my old friend

# Debugging - example (a real one)

- Dual core Debian freeze (RCU error) after ~1h (360'000'000'000 cycles)
  - Impossible to reproduced in simulation (~40 days of runtime @ 100 KHz)
- Hope you get clues
  - Core was still alive on JTAG => OpenOCD
  - PC was in linux “\_raw\_spin\_lock”
  - Forcing it to unlock => Everything back on track
- What it was :
  - CPU D\$ doing bad memory coherency
  - (releasing data permissions once again after being probed out)

# Simulation

- A few open source simulation tools :
  - GHDL, IVVerilog, Verilator, ...
- Based on Verilator (~150-200 KHz on a 4 years old CPU)
- Lock-step checks against Spike / RVLS
- Konata traces

0x80006858: slli a2, a4, 2	F 1 2   Rn 1   Ds 1 2   Is Cm 1 2 3 4 5 6 7
0x8000685c: addi a4, a2, 64	F 1 2   Rn 1   Ds 1 2 3   Is Cm 1 2 3 4 5 6
0x80006860: addi a2, sp, 16	F 1 2   Rn 1   Ds 1 2   Is Cm 1 2 3 4 5 6 7
0x80006864: add a4, a4, a2	F 1 2   Rn 1   Ds 1 2 3   Is Cm 1 2 3 4 5 6
0x80006868: lw a2, 4032(a4)	F 1 2   Rn 1   Ds 1 2 3   Is 1 2 3   Cm 1 2 3
0x8000686c: addiw a2, a2, 1	F 1 2   Rn 1   Ds 1 2 3   Is 1 2 6   Is Cm 1 2 3
0x80006870: sw a2, -64(a4)	F 1 2   Rn 1   Ds 1 2 3   Is 1 2 3 4 5 6 7   Cm
0x80006874: bnez a6, pc + 4048	F 1 2   Rn 1   Ds 1 2   Is 1   Cm 1 2 3 4 5 6 7
0x80006818: lbu a4, 1(a7)	F 1 2   Rn 1   Ds 1 2 3   Is 1 2 3   Cm 1 2 3 4
0x8000681c: addi a2, a7, 1	F 1 2   Rn 1   Ds 1 2   Is Cm 1 2 3 4 5 6 7 8
0x80006820: beq a6, t5, pc + 1148	F 1 2   Rn 1   Ds 1   Is 1   Cm 1 2 3 4 5 6 7 8
0x80006824: addiw t1, a6, 4048	F 1 2   Rn 1   Ds 1 2   Is Cm 1 2 3 4 5 6 7 8
0x80006828: andi t1, t1, 255	F 1 2   Rn 1   Ds 1 2   Is Cm 1 2 3 4 5 6 7 8
0x8000682c: addiw t4, t4, 1	F 1 2   Rn 1   Ds 1 2 3   Is Cm 1 2 3 4 5 6 7
0x80006830: bgeu t3, t1, pc + 628	F 1 2   Rn 1   Ds 1 2   Is 1   Cm 1 2 3 4 5 6 7
0x80006aa4: begz a4, pc + 784	F 1 2   Rn 1   Ds 1 2   Is 1   Cm 1 2 3 4 5 6 7
0x80006aa8: beq a4, t5, pc + 800	F 1 2   Rn 1   Ds 1 2   Is 1   Cm 1 2 3 4 5 6 7
0x80006aac: mv a6, a4	F 1 2   Rn 1   Ds 1   Is Cm 1 2 3 4 5 6 7 8
0x80006ab0: addi a2, a2, 1	F 1 2   Rn 1   Ds 1 2   Is Cm 1 2 3 4 5 6 7 8
0x80006ab4: lbu a4, 0(a2)	F 1 2   Rn 1   Ds 1 2 3   Is 1 2 3   Cm 1 2 3 4
0x80006ab8: beq a6, a3, pc + 44	F 1 2   Rn 1   Ds 1 2   Is 1   Cm 1 2 3 4 5 6 7
0x80006abc: addiw a6, a6, 4048	F 1 2   Rn 1   Ds 1 2 3   Is 1 2 3   Cm 1 2 3 4
0x80006acd: andi a7, a6, 255	F 1 2   Rn 1   Ds 1 2 3   Is 1   Cm 1 2 3 4 5 6 7 8

# 0xAA55

- I hope you had a good time
- If you are looking for open-source project funding :
  - NLnet Foundation (they helped me a lot)

